


Image to LaTeX -- Project Report -- 220704007

1. Problem

We need to convert a 64x256 grayscale image of a math formula into its LaTeX source code. This is an image-to-sequence problem: the input is a fixed-size image 64x256 but the output is a variable-length token sequence (:

before preceding check the file structure:



2. Dataset

All data lives in the `Data_Im2Latx/` folder:

Split	Images	Labels
Train	<code>images_train/</code>	<code>train_formulas.txt</code>
Val	<code>images_val/</code>	<code>validation_formulas.txt</code>

Split	Images	Labels
Test	<code>images_test/</code>	<code>I generated these it lives in the ./test_formulas.txt)</code>

Image filenames are just the formula index (00000.png, 00001.png, ...). Formulas in the text files are one per line, tokens separated by spaces. in simpler way the line number of .txt are the filename of each image.

Preprocessing

- Grayscale, resized to 64x256
- Normalized to [-1, 1] (ToTensor gives [0,1], then Normalize([0.5],[0.5]) maps to [-1,1])
- Light augmentation during training: tiny rotation (1 degree) + shift (2%)
- Tokenized by whitespace; vocabulary built from training set
- Special tokens: `<PAD>` (0), `<SOS>` (1), `<EOS>` (2), `<UNK>` (3)

3. Neural Network Architecture

Based on Deng et al. (2016), "What You Get Is What You See." The model has three main parts connected in a pipeline:

```
Image (1x64x256) --> [CNN Encoder] --> [Row Encoder (biLSTM)] --> [Decoder (LSTM + Attention)] --> LaTeX tokens
```

3.1 CNN Encoder (`ConvEncoder` in `model.py`)

What it does: Extracts visual features from the input image. Each convolutional block learns to detect patterns -- edges in early layers, more complex shapes (symbols, strokes) in later layers.

Structure: Four blocks, each containing:

- `Conv2d(3x3, padding=1)` -- extracts local spatial features
- `BatchNorm2d` -- stabilizes training by normalizing activations
- `ReLU` -- non-linearity so the network can learn complex patterns
- `MaxPool2d(2x2)` -- reduces spatial size by half, keeps strongest features

Channels progression: 1 --> 64 --> 128 --> 256 --> 512

Spatial size after each pool:

Layer	Size
Input	64 x 256
Pool 1	32 x 128
Pool 2	16 x 64
Pool 3	8 x 32
Pool 4	4 x 16

Output: Feature map of shape (B, 512, 4, 16) -- 512 channels, 4 rows, 16 columns.

3.2 Row Encoder (RowEncoder in model.py)

What it does: Reads the CNN feature map left-to-right (and right-to-left) to capture the sequential structure of the formula. A formula like $a + b = c$ has a natural left-to-right order that the CNN alone cannot model.

How it works:

1. Each of the 16 columns is flattened: 512 channels x 4 rows = 2048-dim vector
2. These 16 vectors are fed as a sequence into a bidirectional LSTM
3. Forward LSTM reads left-to-right, backward LSTM reads right-to-left
4. Their outputs are concatenated

Output: Sequence of 16 vectors, each 512-dim (256 forward + 256 backward).

3.3 Decoder (Decoder in model.py)

What it does: Generates LaTeX tokens one at a time, starting from `<SOS>` and stopping at `<EOS>`.

Each step:

1. Embed the previous token (128-dim embedding)
2. Compute attention over encoder outputs --> context vector (what part of the image to look at)
3. Concatenate embedding + context, feed into LSTM (512 hidden units)
4. Concatenate LSTM output + context, project to vocabulary logits
5. Pick the token with highest probability

Initial hidden state: Learned projection of the mean encoder output (not zeros -- this gives the decoder a summary of the whole image to start).

3.4 Attention Mechanism (Attention in model.py)

What it does: Lets the decoder "look at" different parts of the image at each time step. Without attention, the decoder must compress the entire image into one hidden vector. With attention, it can focus on the relevant column (e.g., look at the left side when generating $\frac{1}{2}$, look at the numerator when generating the next symbol).

Type: Bahdanau additive attention.


Formula:

```
score_i = v^T * tanh(W_enc * h_enc_i + W_dec * h_dec)
weights = softmax(scores)           -- how much to attend to each position
context = sum(weights_i * h_enc_i)  -- weighted combination of encoder
outputs
```

Parameters:

- `W_enc`: Linear(512 --> 256) -- projects encoder states
- `W_dec`: Linear(512 --> 256) -- projects decoder hidden state
- `v`: Linear(256 --> 1) -- computes scalar score

3.5 Full Model Summary (`Im2Latex` in `model.py`)



4. Hyperparameters (`config.py`)

Why each value was chosen:

Parameter	Value	Why
Batch size	32	Balances GPU memory usage and gradient noise. Too large (128+) wastes memory and gives less noisy gradients (slower convergence). Too small (4-8) makes training unstable. 32 is a standard default that works well for image-to-seq tasks.
Epochs	40	Enough for the model to converge. The loss/BLEU curves flatten by ~30 epochs. More epochs would risk overfitting.
Optimizer	Adam	Adam adapts the learning rate per-parameter using momentum. Converges faster than plain SGD for sequence models. Standard choice for encoder-decoder architectures.
Learning rate	0.001	Default for Adam. The paper uses similar values. Too high (0.01) causes divergence, too low (0.0001) makes training very slow.
LR schedule	halved every 10 epochs	After initial fast learning, smaller steps help fine-tune. Halving (gamma=0.5) is gentle enough to not kill learning. Every 10 epochs = 4 drops across 40 epochs (0.001 --> 0.0005 --> 0.00025 --> 0.000125).
Gradient clipping	5.0	LSTMs can have exploding gradients (especially early in training). Clipping at 5.0 prevents this without being so aggressive that it slows learning.

Parameter	Value	Why
Teacher forcing	1.0 --> 0.6 (linear decay)	At start, feed ground truth tokens so the model learns faster. Gradually reduce to 60% so the model learns to recover from its own mistakes (exposure bias). Going below 0.5 makes early training unstable.
Max sequence length	200	Longest formula in the dataset. Prevents infinite decoding loops.
CNN filters	[64, 128, 256, 512]	Doubling channels is standard (VGG-style). 512 at the end gives rich features. More would increase memory cost with diminishing returns.
Encoder hidden	256	biLSTM with 256 gives 512-dim output (matching decoder). Larger would add parameters without clear benefit for 16-length sequences.
Embedding dim	128	Standard for vocabularies in the hundreds. Larger (256+) would overparameterize for a small token vocabulary.
Decoder hidden	512	Must be large enough to model long sequences (up to 200 tokens). 512 is the standard in the reference paper.
Attention dim	256	Bottleneck dimension for computing scores. 256 gives enough capacity without being wasteful.
Decoder dropout	0.2	Light regularization to prevent overfitting. 0.2 is standard for decoders. Encoder uses 0.1 (less needed since CNN already regularizes via BN).
Beam width	5	Top-5 candidates during beam search. Diminishing returns beyond 5. Width 1 = greedy. Width 10+ is slow with minimal improvement.

Loss function: Cross-entropy, ignoring PAD tokens (index 0) and the SOS position. This is standard for sequence generation -- it penalizes the model for each wrong token prediction.

5. Files -- What Each One Does

File	Purpose
<code>config.py</code>	All paths + hyperparameters in one place. Change any setting here instead of digging through other files. Controls CNN size, LSTM dims, learning rate, batch size, attention on/off, etc.
<code>vocab.py</code>	Builds the token vocabulary. Reads <code>train_formulas.txt</code> , collects all unique tokens, assigns each an integer ID. Provides <code>encode()</code> (tokens --> IDs) and <code>decode()</code> (IDs --> tokens). Saves/loads to <code>vocab.pkl</code> so it only builds once.

File	Purpose
<code>dataset.py</code>	PyTorch datasets + data loading. <code>FormulaDataset</code> loads train/val image-formula pairs, applies preprocessing (resize, normalize, augment). <code>TestDataset</code> loads test images only (no labels). <code>collate_train</code> pads sequences to equal length in each batch.
<code>model.py</code>	The neural network. Contains all four components: <code>ConvEncoder</code> (CNN), <code>RowEncoder</code> (biLSTM), <code>Attention</code> (Bahdanau), <code>Decoder</code> (LSTM). Also implements <code>greedy()</code> and <code>beam_decode()</code> inference methods.
<code>train.py</code>	Training loop. Runs epochs, computes loss, validates, saves checkpoints, plots curves. Handles teacher forcing decay, learning rate scheduling, gradient clipping. Auto-resumes from checkpoint if one exists.
<code>predict.py</code>	Generates test predictions. Loads the best checkpoint, runs inference on test images using greedy or beam search, optionally applies post-processing. Outputs <code>test_formulas.txt</code> .
<code>postprocess.py</code>	Cleans up generated LaTeX. Fixes common decoder errors: unbalanced braces, empty arguments, double superscripts, repeated tokens (stuttering). Makes output actually compilable.
<code>bleu_score.py</code>	BLEU metric (provided). Computes n-gram overlap between predictions and ground truth.
<code>edit_distance.py</code>	Edit distance metric (provided). Computes normalized Levenshtein distance between predictions and ground truth.
<code>run_all.ipynb</code>	Colab notebook to run the full pipeline step by step: setup, train, evaluate, visualize.



6. Metrics


BLEU (Bilingual Evaluation Understudy): Measures n-gram overlap between predicted and reference token sequences. Ranges from 0 to 1. Higher is better. We use 4-gram BLEU. A score of 0.80 means 80% of predicted n-grams match the reference.

Edit Distance (Normalized Levenshtein): Minimum number of token insertions, deletions, or substitutions needed to convert the prediction into the reference, divided by the length of the longer sequence. We report **accuracy = 1 - normalized_distance**. Higher is better.

7. Results

Training Curves





Visual Comparison

Validation Samples: Ground Truth vs Prediction (green=match, red=diff)

Sample #3400 $S_b = \frac{1}{4\kappa} \left\{ L^2 (R_1^2 + R_2^2) + (R_1)^2 (R_2)^2 \sin^2 \theta \right\}.$	GT: $S_{\{b\}} = \frac{1}{4\kappa} \left\{ L^2 (R_{\{1\}}^2 + R_{\{2\}}^2) \sin^2 \theta \right\}.$ Pred: $S_{\{b\}} = \frac{1}{4\kappa} \left\{ L^2 (R_{\{1\}}^2 + R_{\{2\}}^2) \sin^2 \theta \right\}.$ DIFF
Sample #5422 $[S^z, S^\pm] = \pm S^\pm$	GT: $[S^{\{z\}}, S^{\{\pm\}}] = \pm S^{\{\pm\}}$ Pred: $[S^{\{z\}}, S^{\{\pm\}}] = \pm S^{\{\pm\}}$ MATCH
Sample #8329 $d\Omega_3 = -\sin^2\theta\,dt^2 + \frac{l^2}{r_+^2}(d\theta^2 + \sin^2\theta d\chi^2).$	GT: $d\Omega_{\{3\}} = -\operatorname{operatorname{sin}}^2\theta\,dt^2 + \frac{l^2}{r_+^2}(d\theta^2 + \sin^2\theta d\chi^2).$ Pred: $d\Omega_{\{3\}} = -\operatorname{operatorname{sin}}^2\theta\,dt^2 + \frac{l^2}{r_+^2}(d\theta^2 + \sin^2\theta d\chi^2).$ DIFF
Sample #5160 $J_1 = D, \quad J_2 = -D^3 + 2uD + 2Du$	GT: $J_{\{1\}} = D, \quad J_{\{2\}} = -D^3 + 2uD + 2Du$ Pred: $J_{\{1\}} = D, \quad J_{\{2\}} = -D^3 + 2uD + 2Du$ MATCH


Sample #6023	
$h_{ab}dx^a dx^b = -N_g^2 dt^2 + \sigma_{ij}(dx^i + N^i dt)(dx^j + N^j dt),$	GT: $h_{\{a\}b\}dx^{\{a\}}dx^{\{b\}}=-N_{\{\textcolor{red}{cal}B\}}^{\{2\}}$ Pred: $h_{\{a\}b\}dx^{\{a\}}dx^{\{b\}}=-N_{\{B\}}^{\{2\}}t$ DIFF
Sample #3412	
$\phi(r) = 2f_{\phi}\ln\big(\frac{g}{8\pi f_{\phi}r}\sqrt{\frac{1}{2}-\frac{1}{2N_c}}\big),$	GT: $\phi_{\{r\}}=2f_{\{\phi\}}\operatorname{ln}\big(\frac{g}{8\pi f_{\{\phi\}}}\sqrt{\frac{1}{2}-\frac{1}{2N_c}}\big)$ Pred: $\phi_{\{r\}}=2f_{\{\phi\}}\operatorname{ln}\big(\frac{g}{8\pi f_{\{\phi\}}}\sqrt{\frac{1}{2}-\frac{1}{2N_c}}\big)r$ DIFF
Sample #7913	
$T_{1/2}(u) = (-1)^N \sum_{m,n=1}^2 q^{2m-3} e^{2i(m-n)u} L_{m,n}(u) L_{m,n}^t(u)$	GT: $T_{\{1/2\}}(u) = \{(-1)^{\{N\}}\}\displaystyle\sum_{\{m,n=1\}}^2 q^{2m-3} e^{2i(m-n)u} L_{m,n}(u) L_{m,n}^t(u)$ Pred: $T_{\{1/2\}}(u) = (-1)^{\{N\}}\sum_{\{m,n=1\}}^2 \{$ DIFF
Sample #6676	
$\psi_{\alpha}^{II}(q_3) = \sum_i f_i(q_3) \chi_{\alpha}^i,$	GT: $\psi_{\{\alpha\}}^{\{II\}}(q_{\{3\}}) = \sum_{\{i\}} f_{\{i\}}(q_{\{3\}}) \chi_{\{\alpha\}}^{\{i\}},$ MATCH Pred: $\psi_{\{\alpha\}}^{\{II\}}(q_{\{3\}}) = \sum_{\{i\}} f_{\{i\}}(q_{\{3\}}) \chi_{\{\alpha\}}^{\{i\}},$

Final Validation Numbers

Metric	Value
BLEU (4-gram)	0.831935
Edit Distance Accuracy	0.868808

attention to this point -> higher BLEU is better and lower the Edit Distance but in table its Edit Distance Accuracy so higher is better...

have a look at this (i always say the truth!)



8. Decoding Strategies

Two strategies implemented in model.py:

- **Greedy** (`model.greedy()`): At each step, pick the single highest-probability token. Fast (one forward pass per step), but can miss better sequences.
- **Beam search** (`model.beam_decode()`, width=5): Keep the top-5 partial sequences at each step and expand all of them. Picks the final sequence with the highest cumulative log-probability. Slower (5x more computation) but produces better output because it explores multiple possibilities.

9. Post-Processing (postprocess.py)

The decoder sometimes produces invalid LaTeX. These cleanup steps fix that:

Step	What it fixes	Example
Stutter removal	Decoder repeats tokens 3+ times	<code>\sum \sum \sum \sum --> \sum</code>
Brace balancing	Unmatched <code>{ }</code> <code>()</code> <code>[]</code>	<code>\frac { { a + b } }</code> <code>--> \frac { { a + b } }</code>
Empty argument fix	Commands with empty <code>{ }</code>	<code>\frac { }</code> <code>--> \frac {~}</code>
Double script merge	Two superscripts/subscripts in a row	<code>x ^ {a} ^ {b} --> x ^ { a ^ { b } }</code>
Math wrapping	Missing <code>\$. . .\$</code> delimiters (optional)	<code>x + y --> \$ x + y \$</code>

have a look at this pic:

```
(venv) → CV_220704007 git:(main) ✖ whoami
hasan
(venv) → CV_220704007 git:(main) ✖ python postprocess.py
IN : \frac { x } { }
OUT: $ \frac { x } { } $

IN : \frac { { a + b }
OUT: $ \frac { { a + b }} $

IN : x ^ { 2 } ^ { 3 }
OUT: $ x ^ { 2 ^ { 3 } } $

IN : a _ { i } _ { j }
OUT: $ a _ { i _ { j } } $

IN : \sum \sum \sum \sum _ { i }
OUT: $ \sum \sum _ { i } $

IN : ( [ a + b ]
OUT: $ ( [ a + b ] ) $
```

10. Attention vs No-Attention

IMPORTANT : because of the restriction of Colab (it takes **10** hours to run on **T4**) i cant run the model twice so i have the numbers just with **attention** but the model works with both attention **True** or

False you can change the `config.py` and try this at home:)

Model	BLEU	Edit Dist Accuracy
Without attention	PLS TRAIN THE MODEL YOURSELF	PLS TRAIN THE MODEL YOURSELF
With attention	0.831935	0.868808

How to compare:

1. In `config.py`, set `USE_ATTN = False`
2. Delete or rename the `checkpoints/` folder (so it trains fresh) -> because of the resume logic i have `its for restriction` if you hit train it start from the model that trained 40 epoch you need the rename or set the epoch higher
3. Run training again (cell 6 in notebook)
4. Run evaluation (cells 13-14) and fill the "Without attention" row
5. Set `USE_ATTN = True` back, and fill the "With attention" row from your current results

Expected result: Attention should give better scores because the decoder can look back at specific parts of the image at each step, instead of compressing everything into one hidden vector.

11. References

1. Deng, Yuntian, Anssi Kanervisto, and Alexander M. Rush. "What you get is what you see: A visual markup decompiler." arXiv:1609.04938 (2016).
2. Bahdanau, Dzmitry, Kyunghyun Cho, and Yoshua Bengio. "Neural machine translation by jointly learning to align and translate." ICLR 2015.
3. A Simple Overview of RNN, LSTM and Attention Mechanism [MEDIUM](#)

12. Who am I ?

- **Student** : Hasan Aghaei
- **Professor** : Dr.Adeleh Bitarafan